FlashRNN2: separate recurrence from observation

Generic LSTM / sLSTM | FP16 / BF16 | inference forward | source-observed

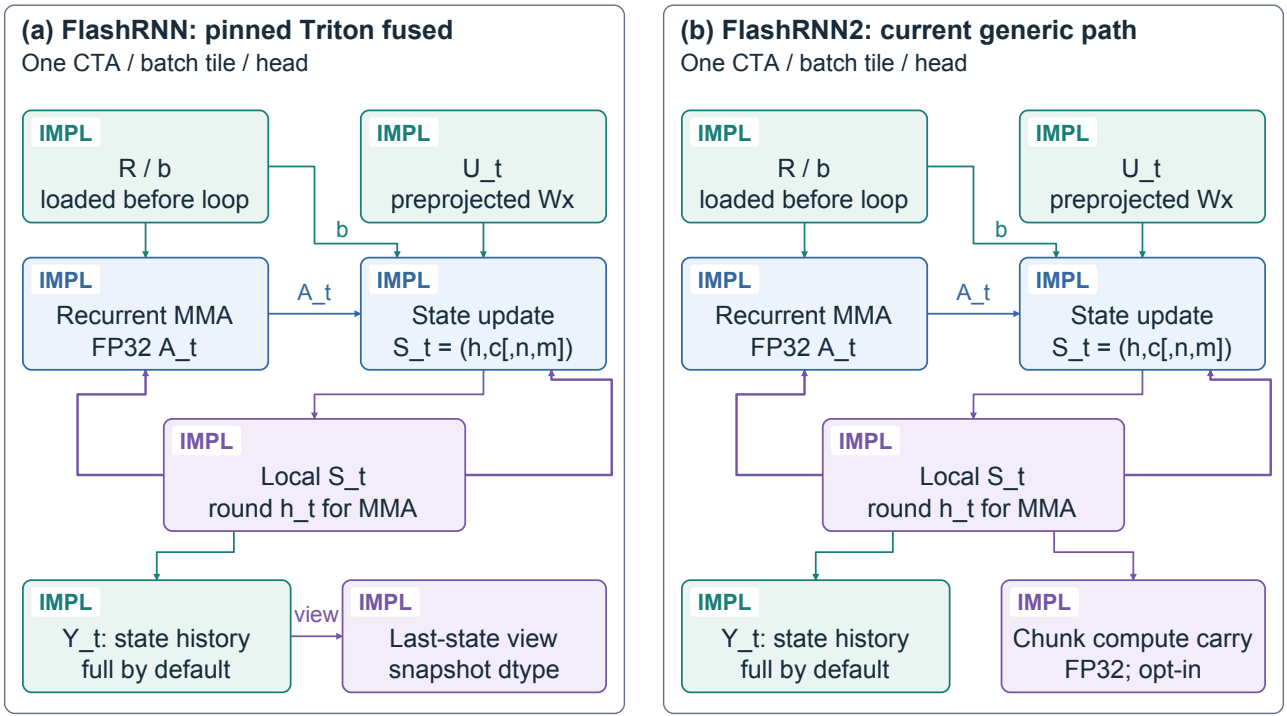

1 Local feedback: inherited
Warp / CTA sync remains

2 Carry + output: separated
Full-call context crosses chunks

3 Overlap: PROPOSED
Layout / spill mapping: TBD

Feedback uses local computation state; required global output writes remain.
Paper CUDA sharding: DOC, grid-sync may apply; selected local baseline is Triton fused.

State / feedback · Compute · Input / output · Sync / constraints · Storage / scope

IMPL: Implemented · **DOC: Documented** · **PROPOSED** · **TBD: Unknown**